

Contents

- ◆ Introduction and fundamentals
- ◆ Introduction to SQL
- ◆ Entity-relationship model
- ◆ Relational model
- ◆ Relational design: normal forms
- ◆ Queries
 - Relational calculus
 - Relational algebra
- ◆ Database implementation
 - Physical structure: fields and records
 - Indexing
 - Simple indexes
 - B trees

Structured Query Language – SQL

- ◆ “Programming” language for DBMS
 - Data definition language: creation of data models (table design)
 - Data manipulation language: insertion, modification, deletion of data
 - Data query language: queries
- ◆ It is run on a DBMS
- ◆ The most widely used standard
 - Created in 1974 (D. D. Chamberlin & R. F. Boyce, IBM)
 - ANSI in 1986, ISO in 1987
 - Core (all DBMS) + packages (optional modules)
- ◆ Versions
 - SQL1 – SQL 86
 - SQL2 – SQL 92
 - SQL3 – SQL 1999, not fully supported by industry (recursion, programming, objects...)
- ◆ Limitations
 - It is not purely relational (e.g. views are *multisets* of tuples)
 - Important divergences between implementations (not directly portable in general, incompleteness, extensions) –one ends up learning SQL variants



Donald D.
Chamberlin

Elements of an SQL database

- ◆ Database = set of tables
- ◆ Table (relation, entity, schema...) =
 - Fixed field structure (schema)
 - Set of records with field values
- ◆ Field (attribute, property, “column”), has a data type
- ◆ Record (tuple, “row”)
- ◆ Primary key
- ◆ Foreign keys

Lexical structure of the language

SQL operations

- ♦ DDL – Table creation, design, modification
- ♦ DML – Record insertion, modification, removal
- ♦ DQL – Query

Lexical structure of SQL

- ♦ Case-insensitive, insignificant whitespace
- ♦ Sentences, expressions, values, data types
- ♦ References
 - Elmasri cap. 8
 - PostgreSQL SQL ref: <http://www.postgresql.org/docs/9.4>

Example: DB for music application

The screenshot displays the Spotify Premium desktop application interface. The top navigation bar includes the Spotify logo, a menu (File, Edit, View, Playback, Help), the text 'Spotify Premium', and window controls. Below this is a search bar and user profile icons. The left sidebar shows navigation options: Stations, Local Files, and a list of Playlists including 'Music For C...', 'Targeted by ...', 'Elena & Pablo', 'Relax & Unw...', 'Reading Sou...', 'Albertucho ...', 'Discover We...', and 'New Pl...'. The main content area features a playlist titled 'Discover Weekly' with a description 'chosen just for you. Updated every Monday, so save your favourites!' and 'Created by: spotifydiscover • 30 songs, 1 hr 57 min'. It includes buttons for 'PLAY', 'FOLLOWING', and 'CREATE SIMILAR PLAYLIST'. Below these is a 'Filter' search bar and an 'Available Offline' toggle. A table lists songs with columns for song title, artist, and date added. The bottom of the interface shows a playback bar with a progress slider (10:56 to 11:09), playback controls (play/pause, previous, next, shuffle, repeat, volume), and a green status bar at the bottom that reads 'You are listening on'.

	SONG	ARTIST	
+	Same To You	Melody Gardot	2 days ago
+	Holding Up - Radio Edit	Jabberwocky, Na Kyung Lee	2 days ago
+	Palmar	Caloncho	2 days ago
+	Sunny	Bryan Adams	2 days ago
+	Jour 1 - Gostan Remix	Louane	2 days ago
+	Quicksand	Caro Emerald	2 days ago

Example: informal description

Music streaming with social network

- ◆ Data types: users, songs, albums, artists...
- ◆ Structures:
 - Users have a name, username, email...
 - Songs have a title, style, duration, date...
 - Artists have name, nationality...
- ◆ Relations:
 - Songs have authors, albums have songs, users have friends, favorite artists, users play songs...
- ◆ Functionalities:
 - Search for a song, play it, browse song data...
 - View / add friends...

Example: table view

Artist

id	name	nationality
0	The Beatles	UK
1	The Rolling Stones	UK
2	David Bowie	UK

Song

id	title	genre	duration	release_date	author
0	Norwegian Wood	Pop	125	1965-03-12	0
1	Here, there and everywhere	Pop	145	1966-08-05	0
2	Jumping jack flash	Pop	225	1968-04-20	1

User

username	name	email
amy	Amelia	amy@gmail.com
jim	James	jim@gmail.com
nick	Nicholas	nick@gmail.com
cate	Catherine	cate@gmail.com

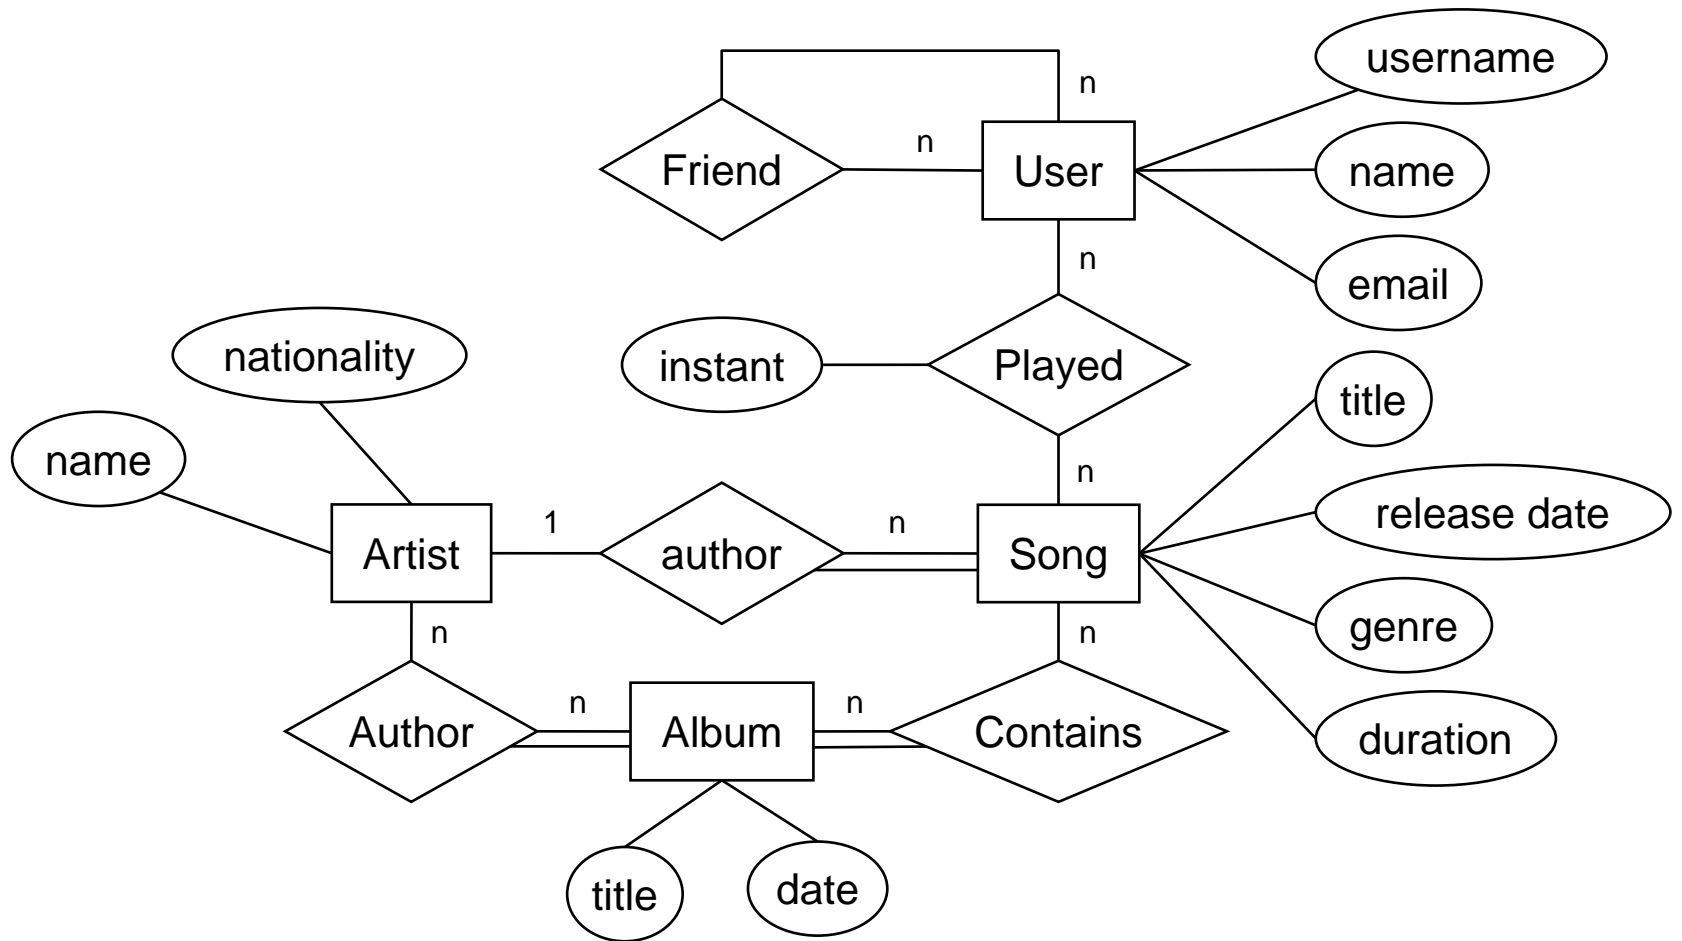
Friend

user1	user2
jim	amy
cate	jim
nick	cate

Played

user	played_song	instant
cate	1	2019-09-09 16:57:54
jim	2	2019-09-12 21:15:30

Example – ER diagram



Motivation

Database for music artist, songs, users, play records, social network

Example queries

Songs of the 60's (show title and genre)

Set of all nationalities of artists in the DB

Songs of artists from the UK

Titles of songs that a user ever listened to

All the friends of a given user

Users with the same name

All users at distance 2 of a given user in the social network

Common friends to two given users

How many times a song has been played

Artists sorted from most to least played

Users with more than two friends


Users sorted by their number of friends

User with most friends

Data definition

```
CREATE TABLE name (  
    field1 type1 [constraints1],  
    field2 type2 [constraints2],  
    ...,  
    [constraints]  
);
```

ALTER commands are useful for slight design modifications of non-empty tables



```
ALTER TABLE name ADD COLUMN field type [constraints];
```

```
ALTER TABLE name ADD constraint;
```

```
ALTER TABLE name DROP COLUMN field;
```

```
DROP TABLE name;
```

```
DROP CONSTRAINT constraint-name;
```

Example

Artist		
id	name	nationality

```
CREATE TABLE Artist (  
    id            int,  
    name          text,  
    nationality    text  
);
```

Example

Artist		
id	name	nationality

```
CREATE TABLE Artist (  
    id            int    PRIMARY KEY,  
    name          text   NOT NULL,  
    nationality    text  
);
```

Example

Song					
id	title	genre	duration	release_date	author

```
CREATE TABLE Song (  
    id            int            PRIMARY KEY,  
    title         text           NOT NULL,  
    genre         text,  
    duration      int,  
    release_date  date,  
    author        int            NOT NULL REFERENCES Artist (id)  
);
```

```
CREATE TABLE User (  
    username    varchar(30) PRIMARY KEY,  
    name        text        NOT NULL,  
    email       text        NOT NULL UNIQUE  
);
```

```
CREATE TABLE Friend (  
    user1       varchar(30) REFERENCES User (username),  
    user2       varchar(30) REFERENCES User (username),  
    PRIMARY KEY (user1, user2)  
);
```

```
CREATE TABLE Played (  
    user        varchar(30),  
    played_song int REFERENCES Song (id),  
    instant     timestamp,  
    PRIMARY KEY (user, played_song, instant)  
);
```

Example: change tables after created

```
ALTER TABLE Played DROP COLUMN instant;
```

```
ALTER TABLE Played ADD instant timestamp; /* NULL's */
```

```
ALTER TABLE Played ADD FOREIGN KEY (user)  
REFERENCES User (username);
```

```
ALTER TABLE User ADD PRIMARY KEY (username);
```

Summary

- ♦ **CREATE** TABLE *name* (...);
- ♦ **ALTER** TABLE *name*

{	ADD	{	COLUMN <i>field type</i> [<i>constraints</i>];
			<i>constraint</i> ;
	DROP		COLUMN <i>field</i> ;
- ♦ **DROP**

{	TABLE <i>name</i> ;
	CONSTRAINT <i>constraint-name</i> ;

- Did you say “constraints”?
- Yes, I did.
- What is that?
- Mainly primary and foreign keys.
We’ll see that in a few moments,
let us see how to enter data first

Resulting DB state: a set of empty tables

Artist

id	name	nationality

Song

id	title	genre	duration	release_date	author

User

username	name	email

Friend

user1	user2

Played

user	played_song	instant

Let's fill this in!

Data manipulation

INSERT INTO *table* [(*field1*, *field2*, ...)]

VALUES

(*value11*, *value12*, ...),

(*value21*, *value22*, ...),

...

A query can also be nested here:
the returned tuples are inserted

;

UPDATE *table* SET *field1* = *value1*, *field2* = *value2*, ...

/* PostgreSQL: [FROM ...] to form conditions with other tables */

[WHERE ...];

DELETE FROM *table*

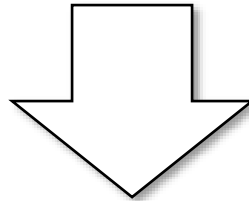
/* PostgreSQL: [USING *table1*, *table2*, ...] for conditions with other tables */

[WHERE ...];

TRUNCATE TABLE *table*;

Example: insert rows

```
INSERT INTO Artist VALUES (0, 'The Beatles', 'UK');
```



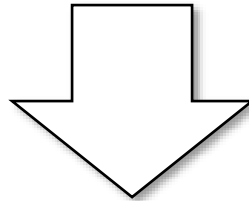
Artist

id	name	nationality
0	The Beatles	UK

Example: insert rows

```
INSERT INTO Artist VALUES (0, 'The Beatles', 'UK');
```

```
INSERT INTO Artist VALUES (1, 'The Rolling Stones', 'UK');
```



Artist

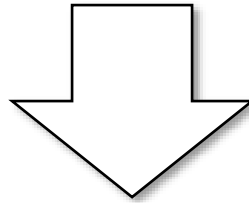
id	name	nationality
0	The Beatles	UK
1	The Rolling Stones	UK

Example: insert **incomplete** rows

INSERT INTO Artist VALUES (0, 'The Beatles', 'UK');

INSERT INTO Artist VALUES (1, 'The Rolling Stones', 'UK');

INSERT INTO **Artist (id, name)** VALUES (2, 'David Bowie');



Artist

id	name	nationality
0	The Beatles	UK
1	The Rolling Stones	UK
2	David Bowie	NULL

Example: insert in the other tables

INSERT INTO Song VALUES

```
(0, 'Norwegian wood', 'Pop', 125, '1965-03-12', 0),  
(1, 'Here, there and everywhere', 'Pop', 145, '1966-08-05', 0),  
(2, 'Jumping jack flash', 'Pop', 225, '1968-04-20', 1);
```

INSERT INTO User VALUES

```
('amy', 'Amelia', 'amy@gmail.com'),  
( 'jim', 'James', 'jim@gmail.com'),  
( 'nick', 'Nicholas', 'nick@gmail.com'),  
( 'cate', 'Catherine', 'cate@gmail.com');
```

INSERT INTO Friend VALUES

```
('jim', 'amy'),  
( 'cate', 'jim'),  
( 'nick', 'cate');
```

INSERT INTO Played VALUES

```
('cate', 2, '2019-09-09 16:57:54'),  
( 'jim', 3, '2019-09-12 21:15:30');
```

Artist

id	name	nationality
0	The Beatles	UK
1	The Rolling Stones	UK
2	David Bowie	NULL

Resulting DB state

Song

id	title	genre	duration	release_date	author
0	Norwegian Wood	Pop	125	1965-03-12	0
1	Here, there and everywhere	Pop	145	1966-08-05	0
2	Jumping jack flash	Pop	225	1968-04-20	1

User

username	name	email
amy	Amelia	amy@gmail.com
jim	James	jim@gmail.com
nick	Nicholas	nick@gmail.com
cate	Catherine	cate@gmail.com

Friend

user1	user2
jim	amy
cate	jim
nick	cate

Played

user	played_song	instant
cate	1	2019-09-09 16:57:54
jim	2	2019-09-12 21:15:30

Example: change cell values

```
UPDATE Artist SET nationality = 'UK'  
WHERE name = 'David Bowie';
```

```
UPDATE Album SET price = price * 1.2;  
/* Fictional field, for illustrative purpose */
```

```
DELETE FROM Played WHERE instant < '2000-01-01 00:00:00';
```

- ◆ What literal values can I use?
- ◆ What data types do I have?
- ◆ How do I write conditions?
- ◆ How do I write expressions?

- ◆ What literal values can I use?
- ◆ What data types do I have?
- ◆ How do I write conditions?
- ◆ How do I write expressions?

These are common questions
for any “programming language”

Types and expressions

SQL types


character(*n*) \equiv char(*n*), varchar(*n*), text

integer \equiv int, smallint

float, real, double precision

numeric (*precision*, *scale*) \equiv decimal (*precision*, *scale*)

date, time, timestamp


integral digits *fractional digits*
 (0 by default)

Literal values

Strings in between '...'

Numeric values similar e.g. to C

date 'YYYY-MM-DD', time 'HH:MM:SS'

Expressions

Can be used in WHERE, SELECT, SET, DEFAULT, CHECK...

Operators

+ - * / % ^

AND OR NOT

= < > <= >= LIKE ISNULL

string operations: concatenation, like, regular expressions ('%' '_')

Comments

--

/* ... */

Now, about constraints...

Constraints

In a field

NOT NULL

UNIQUE

PRIMARY KEY

REFERENCES *table* (*key*) [(ON DELETE | ON UPDATE)

DEFAULT *value*

With name

CONSTRAINT *name* *constraint*

If omitted, will take
the primary key



(NO ACTION | RESTRICT | CASCADE
| SET NULL | SET DEFAULT)]

In a table

PRIMARY KEY (*field1*, *field2*, ...)

FOREIGN KEY (*field1*, *field2*, ...) REFERENCES *table* (*key1*, *key2*, ...)

UNIQUE (*field1*, *field2*, ...)

CHECK (*expression*)

Primary keys

- ◆ They designate a unique identifier for the rows of a table
- ◆ There can be only one primary key per table, though it can comprise several fields
- ◆ It is very advisable that all tables have their primary key
- ◆ Technically equivalent to UNIQUE plus NOT NULL
- ◆ But they play a different role in indexing (we shall see this later on)
- ◆ Design choice: selection of primary key among several options
 - “Natural” primary key (email, web domain, SSN, ISBN, license plate, etc.)
 - “Artificial” primary key: usually an integer ID, typically self-incremented (serial)

Foreign keys

- ◆ Conceptually comparable to pointers
- ◆ They reference unique fields of another table
- ◆ The reference field is commonly a primary key
 - At least it has to be 'unique'
- ◆ Technically they are not indispensable
- ◆ But they are greatly helpful in enforcing consistency in references!
 1. They throw an error when an attempt is made at inserting or setting a FK field to a value that does not exist in the referenced table
 2. And they automate what should happen when a referenced row is deleted
- ◆ In general it is preferable (more efficient) that FKs have integer type

Keys: in summary...

- ◆ The **primary key** of a table acts as a row **identifier**
 - Plays a similar role to the memory address of data in C
 - Only one per table (albeit not mandatory, always designate one)
 - Can be a natural field or (better) an artificial dedicated field (an integer)
 - Can be composed of several fields
- ◆ **Foreign keys** play the part of **pointers** between rows of different (or the same) tables
 - They typically point to a primary key
 - But they can also just point to a unique field
 - Foreign keys can also be composite
- ◆ What is the **advantage** of declaring a **primary key**?
 - Forbid the repetition of values in different rows (by raising an error)
 - Difference to unique not null: a matter of low-level implementation (indexes)
- ◆ What is the **advantage** of declaring **foreign keys**?
 - Enforce that its value appears in the pointed table (raising an error otherwise)
 - React automatically to changes in the pointed primary key

Example: primary and foreign keys

PK

FK

Follows

<u>user1</u>	<u>user2</u>
6	24
73	6
81	73

User

<u>id</u>	name	email
24	Amelia	amy@gmail.com
6	James	jim@gmail.com
81	Nicholas	nick@gmail.com
73	Catherine	cate@gmail.com

Tweet

<u>id</u>	content	author
7	'Congratulations! What an...'	6
48	'Notifications in Twitter are...'	24
35	'I read the first point and it...'	6
92	'Notifications in Twitter are...'	81
50	'I read the first point and it...'	73

Retweet

<u>tweet1</u>	<u>tweet2</u>
92	48
50	35

Example: primary and foreign keys

PK

FK

Follows

<u>user1</u>	<u>user2</u>
6	24
73	6
81	73

User

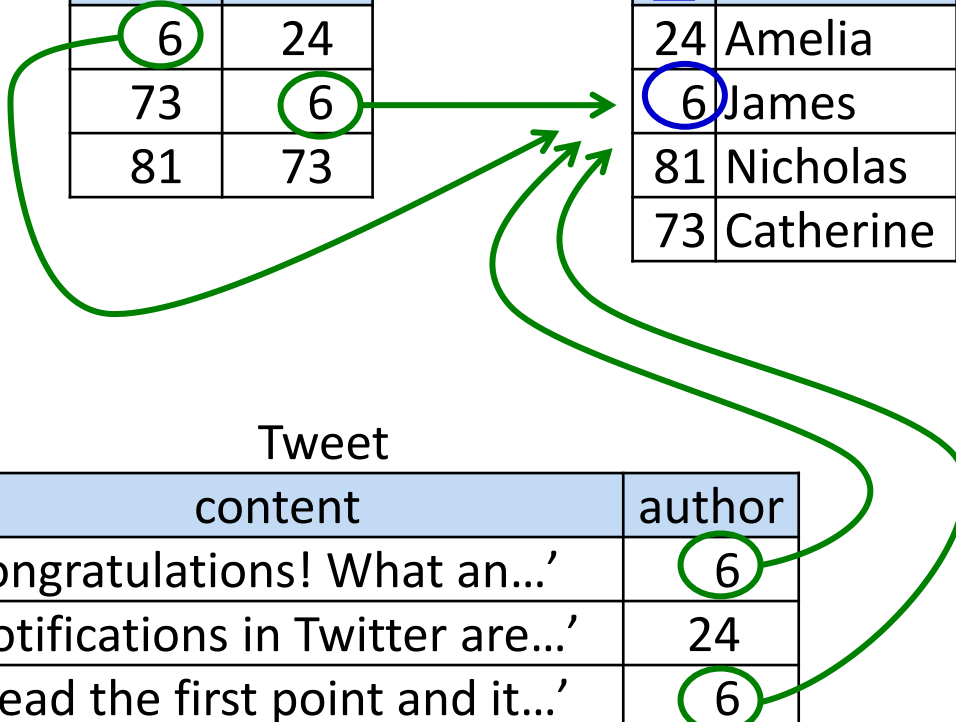
<u>id</u>	name	email
24	Amelia	amy@gmail.com
6	James	jim@gmail.com
81	Nicholas	nick@gmail.com
73	Catherine	cate@gmail.com

Tweet

<u>id</u>	content	author
7	'Congratulations! What an...'	6
48	'Notifications in Twitter are...'	24
35	'I read the first point and it...'	6
92	'Notifications in Twitter are...'	81
50	'I read the first point and it...'	73

Retweet

<u>tweet1</u>	<u>tweet2</u>
92	48
50	35



Example: primary and foreign keys

PK

FK

Follows

<u>user1</u>	<u>user2</u>
6	24
73	6
81	73

User

<u>id</u>	name	email
24	Amelia	amy@gmail.com
6	James	jim@gmail.com
81	Nicholas	nick@gmail.com
73	Catherine	cate@gmail.com

Tweet

<u>id</u>	content	author
7	'Congratulations! What an...'	6
48	'Notifications in Twitter are...'	24
35	'I read the first point and it...'	6
92	'Notifications in Twitter are...'	81
50	'I read the first point and it...'	73

Retweet

<u>tweet1</u>	<u>tweet2</u>
92	48
50	35

Example: primary and foreign keys

PK

FK

Follows

<u>user1</u>	<u>user2</u>
6	24
73	6
81	73

User

<u>id</u>	name	email
24	Amelia	amy@gmail.com
6	James	jim@gmail.com
81	Nicholas	nick@gmail.com
73	Catherine	cate@gmail.com

Tweet

<u>id</u>	content	author
7	'Congratulations! What an...'	6
48	'Notifications in Twitter are...'	24
35	'I read the first point and it...'	6
92	'Notifications in Twitter are...'	81
50	'I read the first point and it...'	73

Retweet

<u>tweet1</u>	<u>tweet2</u>
92	48
50	35

Example: on delete

PK

FK

Follows

<u>user1</u>	<u>user2</u>
6	24
73	6
81	73

User

<u>id</u>	name	email
24	Amelia	amy@gmail.com
6	James	jim@gmail.com
81	Nicholas	nick@gmail.com
73	Catherine	cate@gmail.com

Tweet

<u>id</u>	content	author
7	'Congratulations! What an...'	6
48	'Notifications in Twitter are...'	24
35	'I read the first point and it...'	6
92	'Notifications in Twitter are...'	81
50	'I read the first point and it...'	73

Retweet

<u>tweet1</u>	<u>tweet2</u>
92	48
50	35

Example: on delete cascade

PK

FK

Follows

user1	user2
6	24
73	6
81	73

User

<u>id</u>	name	email
24	Amelia	amy@gmail.com
6	James	jim@gmail.com
81	Nicholas	nick@gmail.com
73	Catherine	cate@gmail.com

Tweet

<u>id</u>	content	author
7	'Congratulations! What an...'	6
48	'Notifications in Twitter are...'	24
35	'I read the first point and it...'	6
92	'Notifications in Twitter are...'	81
50	'I read the first point and it...'	73

Retweet

tweet1	tweet2
92	48
50	35

Cascade

Cascade

Example: on delete cascade

PK

FK

Follows

user1	user2
6	24
73	6
81	73

User

<u>id</u>	name	email
24	Amelia	amy@gmail.com
6	James	jim@gmail.com
81	Nicholas	nick@gmail.com
73	Catherine	cate@gmail.com

Tweet

<u>id</u>	content	author
7	'Congratulations! What an...'	6
48	'Notifications in Twitter are...'	24
35	'I read the first point and it...'	6
92	'Notifications in Twitter are...'	81
50	'I read the first point and it...'	73

Retweet

tweet1	tweet2
92	48
50	35

Cascade

Cascade

Cascade

Example: on delete

PK

FK

Follows

<u>user1</u>	<u>user2</u>
6	24
73	6
81	73

User

<u>id</u>	name	email
24	Amelia	amy@gmail.com
6	James	jim@gmail.com
81	Nicholas	nick@gmail.com
73	Catherine	cate@gmail.com

Tweet

<u>id</u>	content	author
7	'Congratulations! What an...'	6
48	'Notifications in Twitter are...'	24
35	'I read the first point and it...'	6
92	'Notifications in Twitter are...'	81
50	'I read the first point and it...'	73

Retweet

<u>tweet1</u>	<u>tweet2</u>
92	48
50	35

Example: on delete set null

PK

FK

Follows

<u>user1</u>	<u>user2</u>
6	24
73	6
81	73

Set NULL

User

<u>id</u>	name	email
24	Amelia	amy@gmail.com
6	James	jim@gmail.com
81	Nicholas	nick@gmail.com
73	Catherine	cate@gmail.com

Set NULL

Tweet

<u>id</u>	content	author
7	'Congratulations! What an...'	6
48	'Notifications in Twitter are...'	24
35	'I read the first point and it...'	6
92	'Notifications in Twitter are...'	81
50	'I read the first point and it...'	73

Retweet

<u>tweet1</u>	<u>tweet2</u>
92	48
50	35

Example: on delete set null

PK

FK

Follows

<u>user1</u>	<u>user2</u>
NULL	24
73	NULL
81	73

Set NULL

User

<u>id</u>	name	email
24	Amelia	amy@gmail.com
6	James	jim@gmail.com
81	Nicholas	nick@gmail.com
73	Catherine	cate@gmail.com

Set NULL

Tweet

<u>id</u>	content	author
7	'Congratulations! What an...'	NULL
48	'Notifications in Twitter are...'	24
35	'I read the first point and it...'	NULL
92	'Notifications in Twitter are...'	81
50	'I read the first point and it...'	73

Retweet

<u>tweet1</u>	<u>tweet2</u>
92	48
50	35

It is not affected

Example: on update

PK

FK

Follows

<u>user1</u>	<u>user2</u>
6	24
73	6
81	73

User

<u>id</u>	name	email
24	Amelia	amy@gmail.com
6	James	jim@gmail.com
81	Nicholas	nick@gmail.com
73	Catherine	cate@gmail.com

Tweet

<u>id</u>	content	author
7	'Congratulations! What an...'	6
48	'Notifications in Twitter are...'	24
35	'I read the first point and it...'	6
92	'Notifications in Twitter are...'	81
50	'I read the first point and it...'	73

Retweet

<u>tweet1</u>	<u>tweet2</u>
92	48
50	35

Example: on update cascade

PK

FK

Follows

<u>user1</u>	<u>user2</u>
6	24
73	6
81	73

User

<u>id</u>	name	email
24	Amelia	amy@gmail.com
6	James	jim@gmail.com
81	Nicholas	nick@gmail.com
73	Catherine	cate@gmail.com

Tweet

<u>id</u>	content	author
7	'Congratulations! What an...'	6
48	'Notifications in Twitter are...'	24
35	'I read the first point and it...'	6
92	'Notifications in Twitter are...'	81
50	'I read the first point and it...'	73

Cascade

Retweet

<u>tweet1</u>	<u>tweet2</u>
92	48
50	35

Example: on update cascade

PK

FK

Follows

<u>user1</u>	<u>user2</u>
15	24
73	15
81	73

User

<u>id</u>	name	email
24	Amelia	amy@gmail.com
15	James	jim@gmail.com
81	Nicholas	nick@gmail.com
73	Catherine	cate@gmail.com

Tweet

<u>id</u>	content	author
7	'Congratulations! What an...'	15
48	'Notifications in Twitter are...'	24
35	'I read the first point and it...'	15
92	'Notifications in Twitter are...'	81
50	'I read the first point and it...'	73

Cascade

Retweet

<u>tweet1</u>	<u>tweet2</u>
92	48
50	35

It is not affected

What if references are deleted?

PK

FK

Follows

user1	user2
6	24
73	6
81	73

User

id	name	email
24	Amelia	amy@gmail.com
6	James	jim@gmail.com
81	Nicholas	nick@gmail.com
73	Catherine	cate@gmail.com

Tweet

id	content	author
7	'Congratulations! What an...'	6
48	'Notifications in Twitter are...'	24
35	'I read the first point and it...'	6
92	'Notifications in Twitter are...'	81
50	'I read the first point and it...'	73

Retweet

tweet1	tweet2
92	48
50	35

What if references are deleted?

PK

FK

Follows

user1	user2
6	24
73	6
81	73

User

id	name	email
24	Amelia	amy@gmail.com
6	James	jim@gmail.com
81	Nicholas	nick@gmail.com
73	Catherine	cate@gmail.com

No consequence

Tweet

id	content	author
7	'Congratulations! What an...'	6
48	'Notifications in Twitter are...'	24
35	'I read the first point and it...'	6
92	'Notifications in Twitter are...'	81
50	'I read the first point and it...'	73

Retweet

tweet1	tweet2
92	48
50	35

Motivation

Database for music artist, songs, users, play records, social network

Example queries

Songs of the 60's (show title and genre)

Set of all nationalities of artists in the DB

Songs of artists from the UK

Titles of songs that a user ever listened to

All the friends of a given user

Users with the same name

All users at distance 2 of a given user in the social network

Common friends to two given users

How many times a song has been played

Artists sorted from most to least played

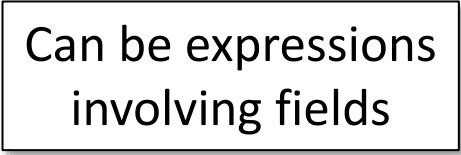
Users with more than two friends

Users sorted by their number of friends

User with most friends

Data query

SELECT [DISTINCT] *fields* FROM *tables*
[WHERE *condition*];



Can be expressions
involving fields

Examples:

```
SELECT title, genre FROM Song  
WHERE release_date > '1959-12-31' AND release_date < '1970-01-01';
```

```
SELECT DISTINCT nationality FROM Artist;
```

```
/* Several tables */
```

```
SELECT * FROM Song, Artist
```

```
WHERE Song.author = Artist.id AND Artist.nationality = 'UK';
```

```
/* Expressions */
```

```
SELECT sid, theory * 0.6 + labs * 0.4 FROM Grades;
```

Motivation

Database for music artist, songs, users, play records, social network

Example queries

Songs of the 60's (show title and genre)

Set of all nationalities of artists in the DB

Songs of artists from the UK

Titles of songs that a user ever listened to

All the friends of a given user

Users with the same name

All users at distance 2 of a given user in the social network

Common friends to two given users

How many times a song has been played

Artists sorted from most to least played

Users with more than two friends

Users sorted by their number of friends

User with most friends

Join

SELECT *fields*

FROM *table1* **JOIN** *table2* **ON** *condition*

[WHERE *condition*];

Typically (though not only) with foreign keys: ON *foreign* = *primary*

More efficient (?) than cartesian product (i.e. just aggregating tables)

Examples:

Semantically
equivalent {

```
SELECT title FROM Song, Played
WHERE user = 'amy' AND played_song = id;

SELECT title FROM Song JOIN Played ON played_song = id
WHERE user = 'amy';
```

```
SELECT * FROM Friend JOIN User
ON (user1 = username OR user2 = username)
WHERE name = 'Catherine';
```

Types of join

- ◆ INNER By default (no need to indicate it)
- ◆ LEFT | RIGHT | FULL Rows that do not fulfil the condition are also included (incompatible with INNER), especially useful in certain queries with aggregative operations
- ◆ NATURAL The condition is of equality between fields with the same name

table1 JOIN *table2* USING (*fields*) is equivalent to a natural join restricted to the fields indicated by 'using'.

Types of join (cont)

Example

```
CREATE TABLE Student (  
    sid VARCHAR(12) PRIMARY KEY, name text);
```

```
CREATE TABLE Course (  
    cid NUMERIC PRIMARY KEY, name text);
```

```
CREATE TABLE Grades (  
    sid VARCHAR(12) REFERENCES Student(sid),  
    cid NUMERIC REFERENCES Course(cid),  
    theory NUMERIC (4,2), labs NUMERIC (4,2),  
    PRIMARY KEY (sid, cid));
```

```
SELECT name, theory FROM Grades NATURAL JOIN Course;
```

```
SELECT name, theory FROM Grades JOIN Course  
ON Grades.cid = Course.cid;
```

Motivation

Database for music artist, songs, users, play records, social network

Example queries

Songs of the 60's (show title and genre)

Set of all nationalities of artists in the DB

Songs of artists from the UK

Titles of songs that a user ever listened to

All the friends of a given user

Users with the same name

All users at distance 2 of a given user in the social network

Common friends to two given users

How many times a song has been played

Artists sorted from most to least played

Users with more than two friends

Users sorted by their number of friends

User with most friends

Alias

SELECT *fields* FROM *table* **AS** *alias* [(*alias-field1*, *alias-field2*, ...)]
[WHERE *condition*];

SELECT *field* **AS** *alias* FROM ...

Example:

SELECT u1.name FROM User **AS** u1, User **AS** u2
WHERE u1.name = u2.name AND u1.username < > u2.username;

SELECT sid, theory * 0.6 + labs * 0.4 **AS** average FROM Grades;

Motivation

Database for music artist, songs, users, play records, social network

Example queries

Songs of the 60's (show title and genre)

Set of all nationalities of artists in the DB

Songs of artists from the UK

Titles of songs that a user ever listened to

All the friends of a given user

Users with the same name

All users at distance 2 of a given user in the social network

Common friends to two given users

How many times a song has been played

Artists sorted from most to least played

Users with more than two friends

Users sorted by their number of friends

User with most friends

Nested queries

SELECT *fields* FROM (SELECT ...) AS *alias* WHERE ...;

SELECT *fields* FROM *table*

WHERE *field1, field2, ...* **IN** (SELECT *field1, field2, ...*);

SELECT *fields* FROM *table*

WHERE *field comparison* (**SOME** | **ALL**) (SELECT ...);

SELECT *fields* FROM *table*

WHERE [NOT] **EXISTS** (SELECT ...);

SELECT *fields* FROM *table*

WHERE (SELECT ...) **CONTAINS** (SOME | ALL) (SELECT ...);

Nested queries (cont)

Examples:

```
SELECT c1.user1  
FROM Friend AS c1 JOIN  
(SELECT * FROM Friend WHERE Friend.user2 = 'amy')  
AS c2      /* In FROM always with AS */  
ON c1.user2 = c2.user1
```

```
SELECT user1 FROM Friend  
WHERE user2 IN (SELECT user1 FROM Friend WHERE user2 = 'amy')
```

Motivation

Database for music artist, songs, users, play records, social network

Example queries

Songs of the 60's (show title and genre)

Set of all nationalities of artists in the DB

Songs of artists from the UK

Titles of songs that a user ever listened to

All the friends of a given user

Users with the same name

All users at distance 2 of a given user in the social network

Common friends to two given users

How many times a song has been played

Artists sorted from most to least played

Users with more than two friends

Users sorted by their number of friends

User with most friends

Set algebra

<i>query1</i> UNION	<i>query2</i>
<i>query1</i> INTERSECT	<i>query2</i>
<i>query1</i> EXCEPT	<i>query2</i>

- ♦ **Homogeneous tuples**: the sets of tuples need to have the same fields
- ♦ An **implicit DISTINCT** is applied (unless we indicate ALL)
- ♦ Can typically be reformulated as a simple query with AND, OR, NOT in the WHERE condition → the advantage could be of readability in some cases

Example:

```
(SELECT user2 FROM Friend WHERE user1 = 'cate'  
UNION  
SELECT user1 FROM Friend WHERE user2 = 'cate')  
INTERSECT  
(SELECT user2 FROM Friend WHERE user1 = 'amy'  
UNION  
SELECT user1 FROM Friend WHERE user2 = 'amy')
```

Motivation

Database for music artist, songs, users, play records, social network

Example queries

Songs of the 60's (show title and genre)

Set of all nationalities of artists in the DB

Songs of artists from the UK

Titles of songs that a user ever listened to

All the friends of a given user

Users with the same name

All users at distance 2 of a given user in the social network

Common friends to two given users

How many times a song has been played

Artists sorted from most to least played

Users with more than two friends

Users sorted by their number of friends

User with most friends

Order and aggregation

SELECT COUNT ([DISTINCT] *field*) FROM *table* ...
[**GROUP BY** *field1*, *field2*, ... [**HAVING** *condition*]];

Post-aggregation
filter



SELECT SUM | MAX | MIN | AVG (*field*) FROM *table* ...
[**GROUP BY** *field1*, *field2*, ...];

SELECT ...

[**ORDER BY** *field1*, *field2*, ... [DESC]];

Why?



Useful combined
with LIMIT *n*

In general when GROUP BY is used, only
those fields can be used in the SELECT
clause (though some DBMS tolerate it)

Order and aggregation: examples

```
SELECT COUNT (*) FROM Played JOIN Song ON played_song = id  
WHERE title = 'Norwegian Wood';
```

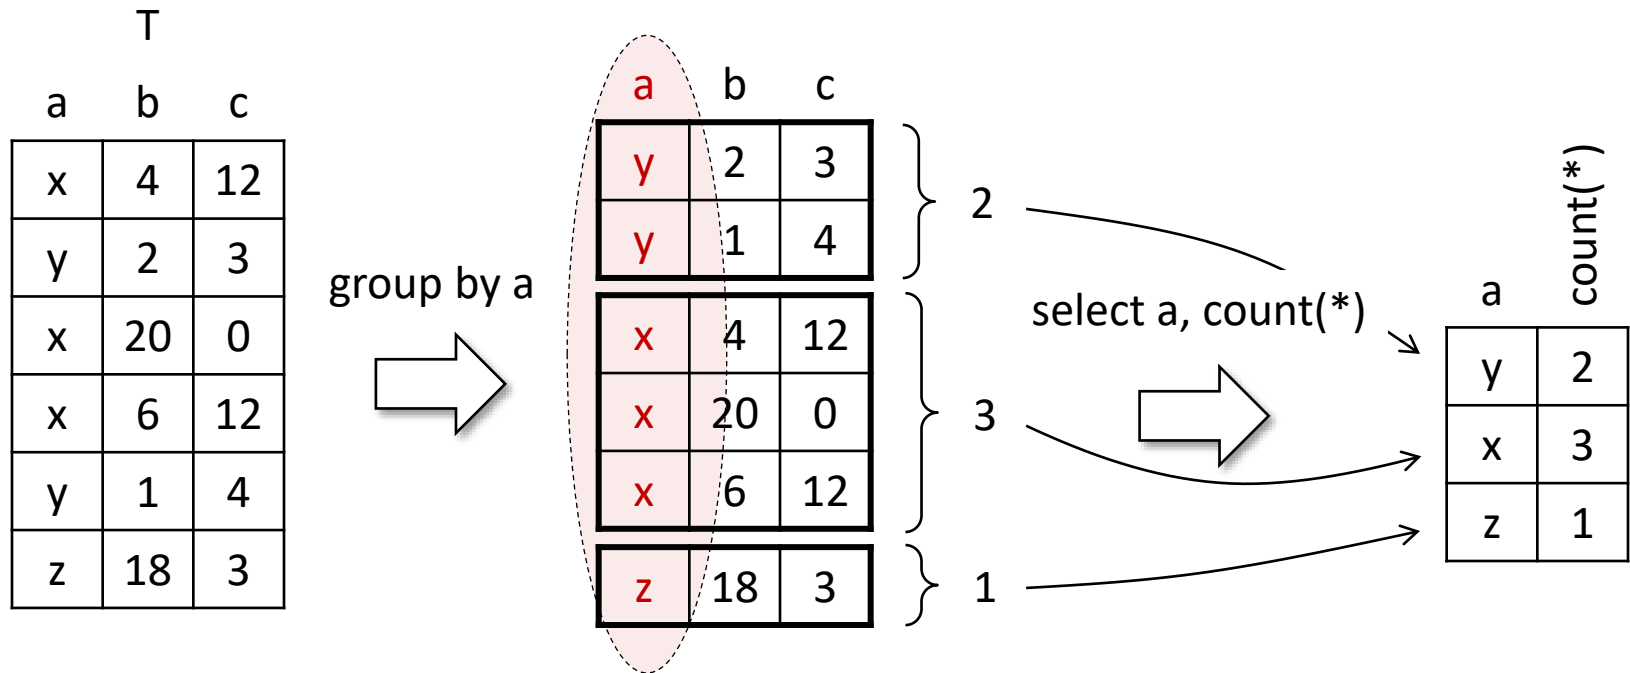
```
SELECT author, COUNT (*) FROM Played JOIN Song ON played_song = id  
GROUP BY author;
```

```
SELECT author, COUNT (*) AS n FROM Played JOIN Song ON played_song = id  
GROUP BY author  
ORDER BY n  
DESC  
LIMIT 1
```

```
SELECT author FROM Played JOIN Song ON played_song = id  
GROUP BY author  
HAVING COUNT (*) > 100
```


Order and aggregation: example

SELECT a, count(*) FROM T GROUP BY a



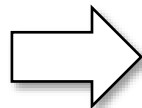
Order and aggregation: example

SELECT a, count(*), **sum(b)** FROM T GROUP BY a

T

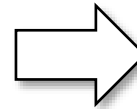
a	b	c
x	4	12
y	2	3
x	20	0
x	6	12
y	1	4
z	18	3

group by a



a	b	c
y	2	3
y	1	4
x	4	12
x	20	0
x	6	12
z	18	3

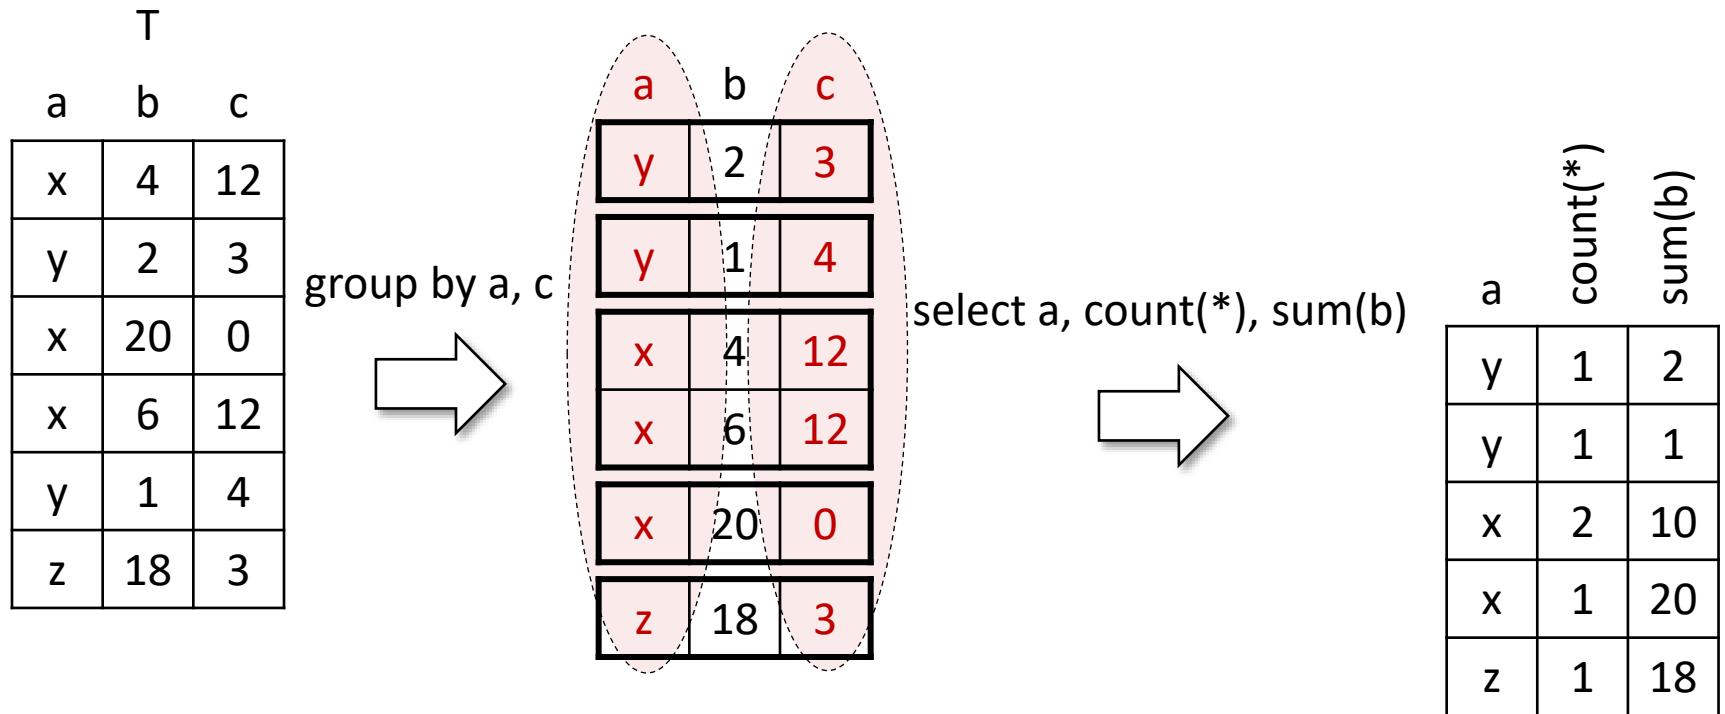
select a, count(*), sum(b)



a	count(*)	sum(b)
y	2	3
x	3	30
z	1	18

Order and aggregation: example

SELECT a, count(*), sum(b) FROM T GROUP BY a, c

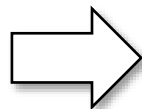


Order and aggregation: example

SELECT a, count(*), sum(b) as s FROM T GROUP BY a, c
HAVING s < 15

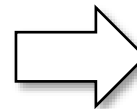
T		
a	b	c
x	4	12
y	2	3
x	20	0
x	6	12
y	1	4
z	18	3

group by a, c



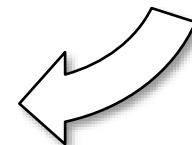
a	b	c
y	2	3
y	1	4
x	4	12
x	6	12
x	20	0
z	18	3

select a, count(*), sum(b)



a	count(*)	s
y	1	2
y	1	1
x	2	10

a	count(*)	s
y	1	2
y	1	1
x	2	10
x	1	20
z	1	18

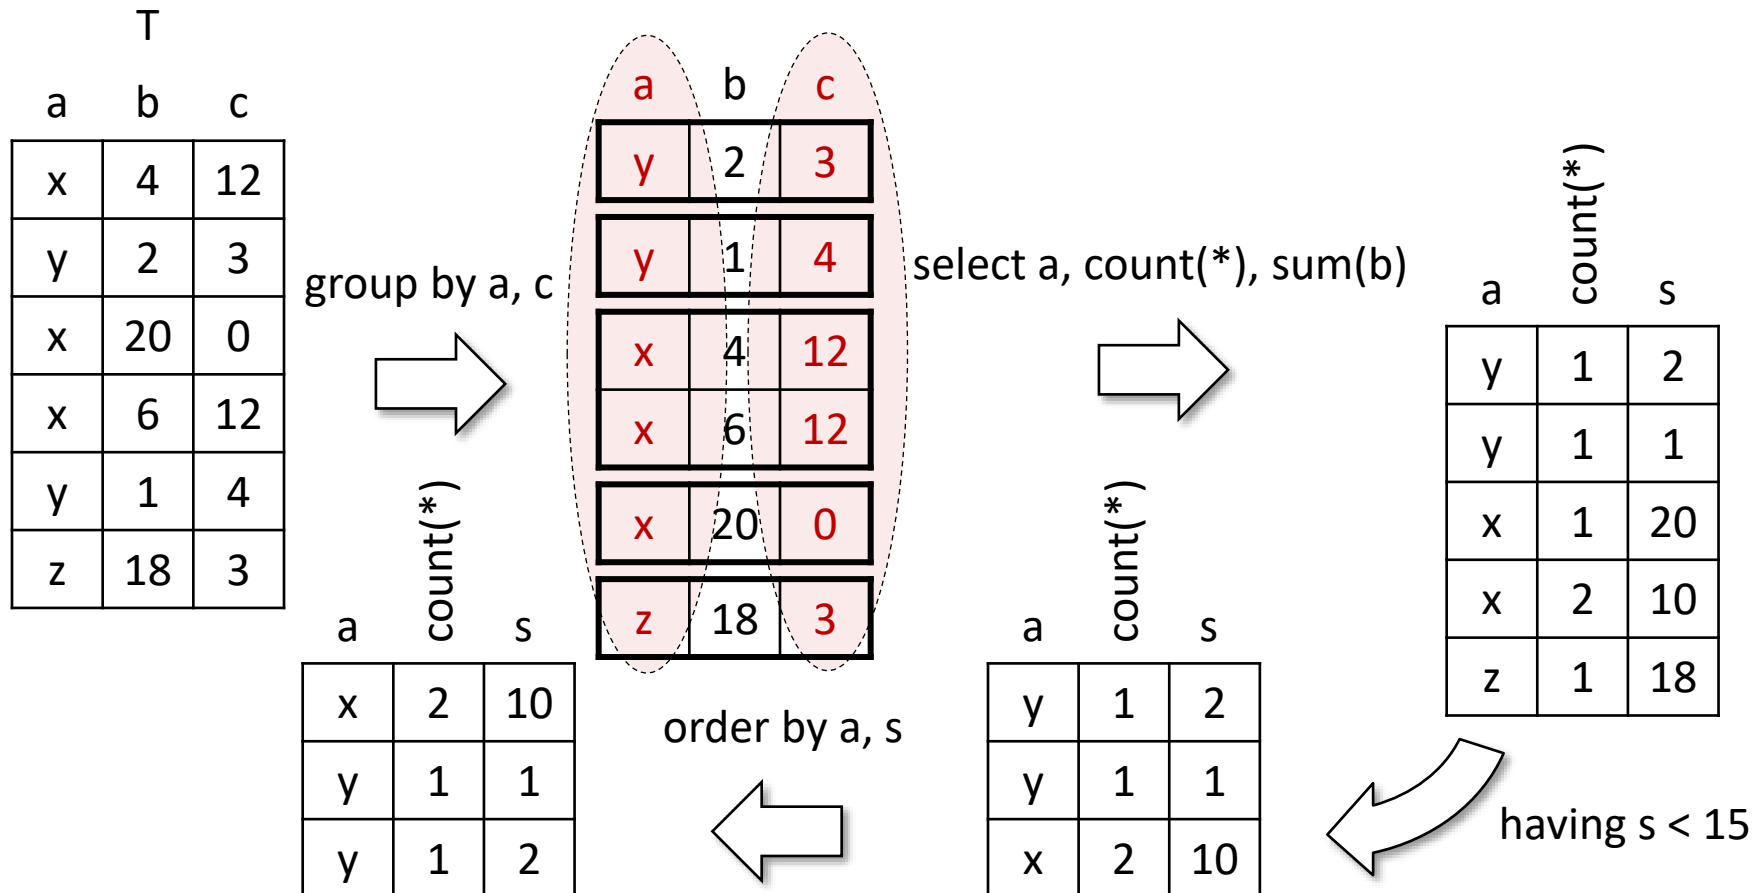


having s < 15

Order and aggregation: example

SELECT a, count(*), sum(b) as s FROM T GROUP BY a, c
HAVING s < 15

ORDER BY a, s



Views

CREATE VIEW *name* **AS** SELECT...;

A way to assign a name to a query

Equivalent to nested query, but...

- Useful to reuse nested queries and improve readability of complex queries
- Views are always up to date when the data change
- Can be configured to be stored on disk (to save re-execution cost)

Examples:

```
CREATE VIEW UserFriends AS
```

```
SELECT u1.username, u2.name FROM User AS u1, User AS u2
```

```
WHERE (u1.username, u2.username) IN
```

```
((SELECT user1, user2 FROM Friend)
```

```
UNION (SELECT user2, user1 FROM Friend));
```

```
SELECT name FROM UserFriends WHERE username = 'jim';
```

Other elements of SQL...

- ◆ Schemas
 - To define namespaces for tables, similar to e.g. Java packages
- ◆ Domains
 - Data types defined by properties and conditions over a primitive data type (e.g. a string with a certain format defined by a regular expression)
- ◆ Triggers
 - Run a procedure when certain update actions (insert, update, delete) are produced in a table
- ◆ Asserts
 - Taylor-made checks over several rows and/or several tables
- ◆ Transactions
 - Establish sequences of actions and queries that must always be completed or cancelled together as a block
 - Also allow to synchronize (block) concurrent operations
- ◆ And many more basic functionalities supported by each DBMS, extending the SQL standard...

Practical informal comments

- ◆ When **several tables** are included in **in the from clause**, two types of conditions will be usually defined in the where clause:
 - The ones expressed in the ‘natural’ query
 - The ones that connect tables to each other (not expressed in the query)
 - Typically foreign key = primary key
 - Each table is connected to some other, in such a way that all tables are connected
 - This is not mandatory, but it is usual for a query to make sense
- ◆ The **alias** are used mainly...
 - On tables: when the same table is involved several times in the same query
 - On fields: to use fields defined by operations (typically in nested queries, or with aggregation)
- ◆ In an initial version of a query it may not be worthwhile to use **join**
 - Let the SQL engine optimize the order of operations
 - If (or when) we have a crisp idea of the optimal order, then go for **joins**
 - External joins can also be useful in certain queries with aggregation

Practical informal comments (cont)

- ♦ **Nested queries** are typically useful...
 - To facilitate complex operations using e.g. not exists, not in
 - To concatenate aggregations, e.g. “Average [avg] number [count] of followers”
 - As views, to ease up complex queries and/or reuse subqueries
 - In other cases it is very often possible to write queries without nesting
- ♦ Queries can generally be written in different equivalent ways
 - Some ways are usually more efficient than others
 - In particular, queries with IN can be generally rewritten with EXISTS and vice-versa, with EXCEPT when they are negated, etc.
 - And any query with OR, AND and NOT in the where clause can be rewritten as a UNION, INTERSECT, EXCEPT
- ♦ The order in which the parts of a query are executed is:
 - FROM, WHERE, GROUP BY, HAVING, SELECT, ORDER BY